

Geometric Computer Vision: Final Project

This notebook contains a revisited implementation of the paper: **Shape Non-rigid Kinematics (SNK): A Zero-Shot Method for Non-Rigid Shape Matching via Unsupervised Functional Map Regularized Reconstruction** by Attaiki and Ovsjanikov (2024).

Imports

```
from pyFM.functional import FunctionalMapping
from model_SURFMNet import SURFMNet
import numpy as np
import trimesh
from pyFM.mesh import TriMesh
import torch
from torch_geometric.data import Batch
from diffusion_net import DiffusionData, DiffusionOperatorsTransform, DiffusionNet
from prism_decoder import PrismDecoder
import torch.optim as optim
from tqdm import tqdm
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d.art3d import Poly3DCollection
from matplotlib import cm
from loss import PrismRegularizationLoss
```

Load meshes

Let's first load the meshes of two dogs modified (non-rigid deformation using Blender):

```
# load some meshes
obj1 = r"./Samples/dog_small.obj"
obj2 = r"./Samples/dog_small_3.obj"
mesh1, mesh2 = TriMesh(obj1), TriMesh(obj2)
```

Compute correspondance from mesh2 to mesh1 using functional maps

I chose the pyFM package investigated in HW3 to compute a first correspondance map between

```
process_params = {
    'n_ev': (70, 70), # Number of eigenvalues on source and Target
    'subsample_step': 20, # In order not to use too many descriptors
    'descr_type': 'WKS', # WKS or HKS
}
model = FunctionalMapping(mesh1, mesh2)
```

```

model.preprocess(**process_params, verbose=False)
fit_params = {
    'w_descr': 1e0,
    'w_lap': 1e-2,
    'w_dcomm': 1e-1,
    'w_orient': 0
}
model.fit(**fit_params, verbose=False)
p2p_21 = model.get_p2p(n_jobs=1)

import meshplot as mp

def double_plot_mp(myMesh1, myMesh2, cmap1=None, cmap2=None):
    p = mp.plot(myMesh1.vertlist, myMesh1.facelist, c=cmap1)
    v2_translated = myMesh2.vertlist.copy()
    f2_translated = myMesh2.facelist.copy()
    myMesh2_translated = TriMesh(v2_translated, f2_translated)
    myMesh2_translated.vertlist[:,0] =
myMesh2_translated.vertlist[:,0] + 30
    p.add_mesh(myMesh2_translated.vertlist,
myMesh2_translated.facelist, c=cmap2)

def visu(vertices):
    min_coord, max_coord =
np.min(vertices, axis=0, keepdims=True), np.max(vertices, axis=0, keepdims=
True)
    cmap = (vertices - min_coord) / (max_coord - min_coord)
    return cmap

cmap1 = visu(mesh1.vertlist); cmap2 = cmap1[p2p_21]
double_plot_mp(mesh1, mesh2, cmap1, cmap2)

{"model_id": "3c3b4574ea8945cc9537fb2e0f010ae3", "version_major": 2, "vers
ion_minor": 0}

```

Compute latent vector

Following section 4.2 in the paper, we compute a latent vector (of length d_1) per feature using DiffusionNet++ module, then apply max pooling to get a vector l of length d_1 . Then we aim to reconstruct the first shape by concatenating l to each feature of the second mesh and passing those to the decoder.

```

def compute_l(my_batch, diffusion_net):
    my_batch.x = my_batch.pos.clone()[0]
    output = diffusion_net(my_batch)
    l = output.x.max(dim=0).values
    return l

def reconstruct_s2(my_batch2, v2_t, l, decoder):

```

```

l_expanded = l.unsqueeze(0).repeat(v2_t.shape[0],1)
my_batch2.x = torch.cat((v2_t,l_expanded),dim=1) #.unsqueeze(0)
s3 = decoder(my_batch2)
return s3

```

Encode the first mesh using a DiffusionNet++ model

```

mesh1_diff = trimesh.load(obj1)
v1, f1 = np.array(mesh1.vertices), np.array(mesh1.faces)
v1_t = torch.from_numpy(v1)
f1_t = torch.from_numpy(f1)
data1 = DiffusionData(pos=v1_t, face=f1_t.T)
diffusion_transform = DiffusionOperatorsTransform(n_eig=50) #97
compute the diffusion net operators with 97 eigenvalues
data1 = diffusion_transform(data1)
my_batch = Batch.from_data_list([data1])
my_batch.pos = my_batch.pos.unsqueeze(0)

mesh2_diff = trimesh.load(obj2)
v2, f2 = np.array(mesh2.vertices), np.array(mesh2.faces)
data2 = DiffusionData(pos=torch.from_numpy(v2),
face=torch.from_numpy(f2).T)
diffusion_transform = DiffusionOperatorsTransform(n_eig=50) #97
compute the diffusion net operators with 97 eigenvalues
data2 = diffusion_transform(data2)
my_batch2 = Batch.from_data_list([data2])
my_batch2.pos = my_batch2.pos.unsqueeze(0)
#####
v2_t = torch.Tensor(v2)

my_batch.pos.requires_grad = True
v2_t.requires_grad = True
my_batch2.pos.requires_grad = True

print(my_batch)
print(my_batch2)

DiffusionDataBatch(pos=[1, 541, 3], face=[3, 1078], mass=[541, 541,
nnz=541], evals=[50], evecs=[541, 50], gradX=[541, 541, nnz=3775],
gradY=[541, 541, nnz=3775], L=[541, 541, nnz=3775], batch=[541],
ptr=[2])
DiffusionDataBatch(pos=[1, 487, 3], face=[3, 970], mass=[487, 487,
nnz=487], evals=[50], evecs=[487, 50], gradX=[487, 487, nnz=3397],
gradY=[487, 487, nnz=3397], L=[487, 487, nnz=3397], batch=[487],
ptr=[2])

from torch.sparse import mm
from model_SURFMNet import FunctionalMapNet
diffusion_net = DiffusionNet(3, 42)
my_batch.x = my_batch.pos.clone()[0]

```

```

my_batch2.x = my_batch2.pos.clone()[0]
output1 = diffusion_net(my_batch)
output2 = diffusion_net(my_batch2)
evecs_T_1 = output1.evecs.t() @ output1.mass.to_dense()
evecs_T_2 = output2.evecs.t() @ output2.mass.to_dense()
fm_net = FunctionalMapNet() #SURFMNet
C1, C2 = fm_net(output1.x.unsqueeze(0), output2.x.unsqueeze(0),
evecs_T_1.unsqueeze(0), evecs_T_2.unsqueeze(0))
emb1 = output1.evecs @ evecs1[:, :k1] @ FM_12.T
emb2 = output2.evecs @ C1.squeeze() #evecs1[:, :k1] @ FM_12.T
dist = torch.cdist(emb1, emb2)#, dim=1, p=None)
knn = dist.topk(1, largest=False, dim=0)

def double_plot(myMesh1, myMesh2, cmap1=None, cmap2=None,
title1='Mesh1',title2='Mesh2'):
    fig = plt.figure(figsize=plt.figaspect(0.5))
    ax1 = fig.add_subplot(1, 2, 1, projection='3d')
    ax1.plot_trisurf(myMesh1.vertlist[:, 0], myMesh1.vertlist[:, 1],
myMesh1.vertlist[:, 2],
                    triangles=myMesh1.facelist, cmap='viridis',
facecolors=cmap1)
    ax1.set_title(title1)
    ax2 = fig.add_subplot(1, 2, 2, projection='3d')
    ax2.plot_trisurf(myMesh2.vertlist[:, 0], myMesh2.vertlist[:, 1],
myMesh2.vertlist[:, 2],
                    triangles=myMesh2.facelist, cmap='viridis',
facecolors=cmap2)
    ax2.set_title(title2)
# cmap1 = visu(v1_remapped)
# cmap2 = visu(s3.features.detach().squeeze(0).numpy())
# double_plot(mesh1, mesh2, cmap1, cmap2, '', '')

def visu_face_colors(s1_faces_raw, s1_faces, s2_faces, s3_faces):
    cmap = cm.get_cmap('viridis', s3_faces.shape[0]) # Use a colormap
with 500 distinct colors
    face_colors = cmap(np.linspace(0, 1, s3_faces.shape[0]))

    average_z = np.mean(s2_faces[:, :, 2], axis=1)
    sorted_indices = np.argsort(average_z)

    sorted_s3_faces = s3_faces[sorted_indices]
    sorted_s2_faces = s2_faces[sorted_indices]
    sorted_s1_faces = s1_faces[sorted_indices]
    fig = plt.figure()
    ax1 = fig.add_subplot(131, projection='3d')
    ax2 = fig.add_subplot(132, projection='3d')
    ax3 = fig.add_subplot(133, projection='3d')
    for i, face in enumerate(s3_faces):
        # poly3d = [[vertices[vert_idx] for vert_idx in face]]

```

```

        poly3d_1 = [sorted_s1_faces[i]]
        poly3d_2 = [sorted_s2_faces[i]]
        poly3d_1_raw = [s1_faces_raw[i]]
        poly3d_3 = [sorted_s3_faces[i]]
        ax1.add_collection3d(Poly3DCollection(poly3d_1,
facecolors=face_colors[i], linewidths=1, edgecolors=None, alpha=.8))
        ax1.add_collection3d(Poly3DCollection(poly3d_1_raw,
facecolors='b', linewidths=1, edgecolors=None, alpha=.1))
        ax2.add_collection3d(Poly3DCollection(poly3d_2,
facecolors=face_colors[i], linewidths=1, edgecolors=None, alpha=.8))
        ax3.add_collection3d(Poly3DCollection(poly3d_3,
facecolors=face_colors[i], linewidths=1, edgecolors=None, alpha=.8))
        ax1.set_xlim([s1_faces[:, :, 0].min(), s1_faces[:, :, 0].max()])
        ax1.set_ylim([s1_faces[:, :, 1].min(), s1_faces[:, :, 1].max()])
        ax1.set_zlim([s1_faces[:, :, 2].min(), s1_faces[:, :, 2].max()])
        ax2.set_xlim([s2_faces[:, :, 0].min(), s2_faces[:, :, 0].max()])
        ax2.set_ylim([s2_faces[:, :, 1].min(), s2_faces[:, :, 1].max()])
        ax2.set_zlim([s2_faces[:, :, 2].min(), s2_faces[:, :, 2].max()])
        ax3.set_xlim([s3_faces[:, :, 0].min(), s3_faces[:, :, 0].max()])
        ax3.set_ylim([s3_faces[:, :, 1].min(), s3_faces[:, :, 1].max()])
        ax3.set_zlim([s3_faces[:, :, 2].min(), s3_faces[:, :, 2].max()])
        plt.show()

def compute_centroid(face):
    return np.mean(face, axis=0)

def find_closest_faces(lf1, lf2):
    centroids_lf1 = np.array([compute_centroid(face) for face in lf1])
    centroids_lf2 = np.array([compute_centroid(face) for face in lf2])

    closest_faces = []
    closest_faces_idx = []

    for centroid in centroids_lf2:
        distances = np.linalg.norm(centroids_lf1 - centroid, axis=1)
        closest_face_idx = np.argmin(distances)
        while closest_face_idx in closest_faces_idx:
            distances[closest_face_idx] = max(distances) + 1
            closest_face_idx = np.argmin(distances)
        closest_faces.append(lf1[closest_face_idx])
        closest_faces_idx.append(closest_face_idx)
    return closest_faces_idx

vls = my_batch.pos.reshape(-1, 3)
fls = my_batch.face.t()
s1_faces_tmp = vls[fls].detach().numpy()

def compute_fmaps(my_batch, my_batch2, diff_net_fts):
    my_batch.x = my_batch.pos.clone()[0]
    my_batch2.x = my_batch2.pos.clone()[0]

```

```

output1 = diff_net_fts(my_batch)
output2 = diff_net_fts(my_batch2)
evecs_T_1 = output1.evecs.t() @ output1.mass.to_dense()
evecs_T_2 = output2.evecs.t() @ output2.mass.to_dense()
fm_net = FunctionalMapNet() #SURFMNet
C1, C2 = fm_net(output1.x.unsqueeze(0), output2.x.unsqueeze(0),
evecs_T_1.unsqueeze(0), evecs_T_2.unsqueeze(0))
emb1 = output1.evecs #evecs1[:, :k1] @ FM_12.T
emb2 = output2.evecs @ C1.squeeze() #evecs1[:, :k1] @ FM_12.T
dist = torch.cdist(emb1, emb2)#, dim=1, p=None)
knn = dist.topk(1, largest=False, dim=0)
p2p_21 = knn.indices.squeeze(0)
return C1, C2, p2p_21

L_SPACE_SIZE = 42
diff_net_fts = DiffusionNet(3, 42)
diffusion_net = DiffusionNet(3, L_SPACE_SIZE)
decoder = PrismDecoder(v1.shape[1]+L_SPACE_SIZE, v1.shape[0])
get_energy_loss = PrismRegularizationLoss(100)

diffusion_net.first_lin.weight.grad

# %debug
# l = compute_l(my_batch)
# s3 = reconstruct_s2(my_batch2, v2_t, l)
optimizer_diffnet = optim.AdamW(diffusion_net.parameters(), lr=0.001,
weight_decay=0.001)
optimizer_diffnet_fts =
optim.AdamW(diff_net_fts.parameters(), lr=0.001, weight_decay=0.001)
optimizer_decoder = optim.AdamW(decoder.parameters(), lr=0.001,
weight_decay=0.001)
for epoch in (pbar:=tqdm(range(10000))):
    l = compute_l(my_batch, diffusion_net)
    l.retain_grad()
    s3 = reconstruct_s2(my_batch2, v2_t, l, decoder)
    s3.features.retain_grad()
    #####
    C1, C2, p2p_21 = compute_fmaps(my_batch, my_batch2, diff_net_fts)
    C1.retain_grad()
    C2.retain_grad()
    # p2p_21.retain_grad()
    #####
    v1_remapped = v1[p2p_21]
    loss_MSE =
torch.nn.functional.mse_loss(torch.Tensor(v1_remapped), s3.features)
    loss_FMAPS_1 = torch.nn.functional.mse_loss(C1@C2, torch.eye(50))
+ torch.nn.functional.mse_loss(C1.transpose(1,2)@C1, torch.eye(50))
    loss_FMAPS_2 = torch.nn.functional.mse_loss(C2@C1, torch.eye(50))
+ torch.nn.functional.mse_loss(C2.transpose(1,2)@C1, torch.eye(50))
    loss_E = get_energy_loss(s3.transformed_prism, s3.rotations,

```

```

s3.pos.reshape(-1, 3), s3.face)
    total_loss = 1e-10*loss_FMAPS_1 + 1e-10*loss_FMAPS_2 + loss_MSE +
    loss_E*10
    pbar.set_description(f"{total_loss.detach():3f}")
    loss_MSE.retain_grad()
    loss_FMAPS_1.retain_grad()
    loss_FMAPS_2.retain_grad()
    loss_E.retain_grad()
    total_loss.backward()
    # import ipdb; ipdb.set_trace()
    # my_batch2.pos = my_batch2.pos.squeeze(0)
    optimizer_diffnet.step()
    optimizer_diffnet.zero_grad()
    optimizer_decoder.step()
    optimizer_decoder.zero_grad()
    optimizer_diffnet_fts.step()
    optimizer_diffnet_fts.zero_grad()
    if epoch%100 == 0:
        print(f'Epoch {epoch}: MSE loss = {loss_MSE}')
        print(f'Epoch {epoch}: Energy loss = {10*loss_E}')
        print(f'Epoch {epoch}: f1 loss = {1e-10*loss_FMAPS_1}')
        print(f'Epoch {epoch}: f2 loss = {1e-10*loss_FMAPS_2}')
        verts = my_batch2.pos.reshape(-1, 3)
        faces = my_batch2.face.t()
        s2_faces = verts[faces].detach().numpy()
        s3_faces = s3.transformed_prism.detach().squeeze(0).numpy()
        # closest_faces_idx = find_closest_faces(s3_faces,
s1_faces_tmp)
        closest_faces_idx =
find_closest_faces(s1_faces_tmp,s3_faces)#, s1_faces_tmp)
        s1_faces = s1_faces_tmp[closest_faces_idx]
        visu_face_colors(s1_faces_tmp, s1_faces, s2_faces, s3_faces)
        plt.show()

```

```

0%|
| 0/10000 [00:00<?, ?it/s]C:\Users\Hadassa-Port\AppData\Local\Temp\
ipykernel_21200\2945219647.py:19: UserWarning: Using a target size
(torch.Size([1, 487, 3])) that is different to the input size
(torch.Size([487, 3])). This will likely lead to incorrect results due
to broadcasting. Please ensure they have the same size.

```

```

    loss_MSE =
torch.nn.functional.mse_loss(torch.Tensor(v1_remapped),s3.features)
C:\Users\Hadassa-Port\AppData\Local\Temp\
ipykernel_21200\2945219647.py:20: UserWarning: Using a target size
(torch.Size([50, 50])) that is different to the input size
(torch.Size([1, 50, 50])). This will likely lead to incorrect results
due to broadcasting. Please ensure they have the same size.

```

```

    loss_FMAPS_1 = torch.nn.functional.mse_loss(C1@C2, torch.eye(50)) +
torch.nn.functional.mse_loss(C1.transpose(1,2)@C1, torch.eye(50))
C:\Users\Hadassa-Port\AppData\Local\Temp\

```




```

ipykernel_21200\2945219647.py:21: UserWarning: Using a target size
(torch.Size([50, 50])) that is different to the input size
(torch.Size([1, 50, 50])). This will likely lead to incorrect results
due to broadcasting. Please ensure they have the same size.
  loss_FMAPS_2 = torch.nn.functional.mse_loss(C2@C1, torch.eye(50)) +
torch.nn.functional.mse_loss(C2.transpose(1,2)@C1, torch.eye(50))
C:\Users\Hadassa-Port\Desktop\hadassa\Toar 2\Semester 1\Geometric
Computer Vision\SNK\loss.py:28: UserWarning: Using torch.cross without
specifying the dim arg is deprecated.
Please either pass the dim explicitly or simply use
torch.linalg.cross.
The default value of dim will change to agree with that of
linalg.cross in a future release. (Triggered internally at ..\aten\
src\ATen\native\Cross.cpp:66.)
  normal = torch.cross(edge1, edge2)
441.944489:   0%|
| 0/10000 [00:00<?, ?it/s]

Epoch 0: MSE loss = 138.4190216064453
Epoch 0: Energy loss = 303.43438720703125
Epoch 0: f1 loss = 0.014822783879935741
Epoch 0: f2 loss = 0.07625257223844528

C:\Users\Hadassa-Port\AppData\Local\Temp\
ipykernel_21200\2975563605.py:3: MatplotlibDeprecationWarning: The
get_cmap function was deprecated in Matplotlib 3.7 and will be removed
two minor releases later. Use ``matplotlib.colormaps[name]`` or
``matplotlib.colormaps.get_cmap(obj)`` instead.
  cmap = cm.get_cmap('viridis', s3_faces.shape[0]) # Use a colormap
with 500 distinct colors

```




```

101.306755:   1%|█
| 100/10000 [02:40<3:41:11, 1.34s/it]


Epoch 100: MSE loss = 84.59368896484375
Epoch 100: Energy loss = 15.418474197387695
Epoch 100: f1 loss = 0.5673588514328003
Epoch 100: f2 loss = 0.7272352576255798

```


389.330383: 2%
| 200/10000 [05:17<3:25:09, 1.26s/it]

Epoch 200: MSE loss = 62.243560791015625
Epoch 200: Energy loss = 14.786492347717285
Epoch 200: f1 loss = 116.22406005859375
Epoch 200: f2 loss = 196.07626342773438




78.396454: 3%
| 300/10000 [07:53<3:23:45, 1.26s/it]

Epoch 300: MSE loss = 62.15915298461914
Epoch 300: Energy loss = 16.153026580810547
Epoch 300: f1 loss = 0.01310831680893898
Epoch 300: f2 loss = 0.07116103917360306




110.424339: 4% | ██████████
 | 400/10000 [10:30<3:19:16, 1.25s/it]

Epoch 400: MSE loss = 85.4991683959961
 Epoch 400: Energy loss = 17.634859085083008
 Epoch 400: f1 loss = 0.0022359283175319433
 Epoch 400: f2 loss = 7.288074493408203




64.594574: 5% | ██████████
 | 500/10000 [13:03<3:17:11, 1.25s/it]

Epoch 500: MSE loss = 54.403465270996094
 Epoch 500: Energy loss = 9.752196311950684
 Epoch 500: f1 loss = 0.37450987100601196
 Epoch 500: f2 loss = 0.06440340727567673




191.702759: 6%|██████████
 | 600/10000 [15:42<3:19:24, 1.27s/it]

Epoch 600: MSE loss = 158.76788330078125
 Epoch 600: Energy loss = 21.617042541503906
 Epoch 600: f1 loss = 8.677755355834961
 Epoch 600: f2 loss = 2.6400654315948486




205.595535: 7%|██████████
 | 700/10000 [18:30<3:09:08, 1.22s/it]

Epoch 700: MSE loss = 192.72781372070312
 Epoch 700: Energy loss = 11.030577659606934
 Epoch 700: f1 loss = 0.14665056765079498
 Epoch 700: f2 loss = 1.6904958486557007




142.739319: 8% | XXXXXXXXXX
 | 800/10000 [21:05<3:09:25, 1.24s/it]

Epoch 800: MSE loss = 131.19796752929688
 Epoch 800: Energy loss = 11.491016387939453
 Epoch 800: f1 loss = 0.03889824077486992
 Epoch 800: f2 loss = 0.011435823515057564




76.922302: 9% | XXXXXXXXXX
 | 900/10000 [23:49<3:42:16, 1.47s/it]

Epoch 900: MSE loss = 64.81265258789062
 Epoch 900: Energy loss = 11.948801040649414
 Epoch 900: f1 loss = 0.005377727095037699
 Epoch 900: f2 loss = 0.1554747372865677




199.078140: 10%|██████████
 | 1000/10000 [26:58<3:32:11, 1.41s/it]

Epoch 1000: MSE loss = 74.0458755493164
 Epoch 1000: Energy loss = 11.746149063110352
 Epoch 1000: f1 loss = 112.19220733642578
 Epoch 1000: f2 loss = 1.0939030647277832




51.597801: 11%|██████████
 | 1100/10000 [30:21<4:00:55, 1.62s/it]

Epoch 1100: MSE loss = 35.986412048339844
 Epoch 1100: Energy loss = 15.169523239135742
 Epoch 1100: f1 loss = 0.029017485678195953
 Epoch 1100: f2 loss = 0.4128505289554596




73.428291: 12%|██████████
 | 1200/10000 [33:41<3:11:45, 1.31s/it]

Epoch 1200: MSE loss = 57.123775482177734
 Epoch 1200: Energy loss = 16.2049503326416
 Epoch 1200: f1 loss = 0.08186553418636322
 Epoch 1200: f2 loss = 0.017696553841233253




46.046829: 13%|██████████
 | 1300/10000 [37:10<3:11:11, 1.32s/it]

Epoch 1300: MSE loss = 31.36875343322754
 Epoch 1300: Energy loss = 14.604534149169922
 Epoch 1300: f1 loss = 0.001092314487323165
 Epoch 1300: f2 loss = 0.07245117425918579




118.135956: 14%| ████████████████████
 | 1400/10000 [41:14<3:11:47, 1.34s/it]

Epoch 1400: MSE loss = 106.15186309814453
 Epoch 1400: Energy loss = 11.93130111694336
 Epoch 1400: f1 loss = 0.04649919271469116
 Epoch 1400: f2 loss = 0.006287424359470606




71.062225: 15%| ████████████████████
 | 1500/10000 [44:36<4:06:58, 1.74s/it]

Epoch 1500: MSE loss = 50.54459762573242
 Epoch 1500: Energy loss = 20.486000061035156
 Epoch 1500: f1 loss = 0.03056538663804531
 Epoch 1500: f2 loss = 0.0010624895803630352



103.634705: 16%| XXXXXXXXXX
 | 1600/10000 [48:37<4:20:22, 1.86s/it]


Epoch 1600: MSE loss = 88.33651733398438
 Epoch 1600: Energy loss = 15.207642555236816
 Epoch 1600: f1 loss = 0.0004653192590922117
 Epoch 1600: f2 loss = 0.09008392691612244



61.896221: 17%| XXXXXXXXXX
 | 1700/10000 [52:01<3:33:07, 1.54s/it]


Epoch 1700: MSE loss = 48.62522506713867
 Epoch 1700: Energy loss = 13.07433032989502
 Epoch 1700: f1 loss = 0.006607468705624342
 Epoch 1700: f2 loss = 0.19005973637104034

```
Epoch 1900: MSE loss = 118.1742172241211
Epoch 1900: Energy loss = 15.890031814575195
Epoch 1900: f1 loss = 0.19845710694789886
Epoch 1900: f2 loss = 0.02494070865213871
```




91.834633: 20%| XXXXXXXXXX
 | 2000/10000 [1:02:09<3:05:55, 1.39s/it]

Epoch 2000: MSE loss = 66.97698974609375
 Epoch 2000: Energy loss = 13.020589828491211
 Epoch 2000: f1 loss = 11.77174186706543
 Epoch 2000: f2 loss = 0.06530854105949402




194.612015: 21%| XXXXXXXXXX
 | 2100/10000 [1:05:34<2:55:52, 1.34s/it]

Epoch 2100: MSE loss = 44.03076934814453
 Epoch 2100: Energy loss = 13.104503631591797
 Epoch 2100: f1 loss = 107.11314392089844
 Epoch 2100: f2 loss = 30.36359977722168



92.570679: 24% | XXXXXXXXXX
 | 2400/10000 [1:15:31<2:49:36, 1.34s/it]


Epoch 2400: MSE loss = 76.92635345458984
 Epoch 2400: Energy loss = 15.631635665893555
 Epoch 2400: f1 loss = 0.0034903977066278458
 Epoch 2400: f2 loss = 0.009200639091432095



71.552452: 25% | XXXXXXXXXX
 | 2500/10000 [1:18:57<2:57:07, 1.42s/it]


Epoch 2500: MSE loss = 47.19223403930664
 Epoch 2500: Energy loss = 14.415288925170898
 Epoch 2500: f1 loss = 9.758109092712402
 Epoch 2500: f2 loss = 0.1868216097354889

```
Epoch 2700: MSE loss = 130.90032958984375
Epoch 2700: Energy loss = 12.551244735717773
Epoch 2700: f1 loss = 1.1347119808197021
Epoch 2700: f2 loss = 1.0387150049209595
```


```
69.109978: 34%|███████████  
| 3400/10000 [1:49:43<2:49:32, 1.54s/it]
```

```
Epoch 3400: MSE loss = 57.740577697753906
Epoch 3400: Energy loss = 11.366235733032227
Epoch 3400: f1 loss = 0.002429555868729949
Epoch 3400: f2 loss = 0.0007368364022113383
```



```
134.334839: 35%|
| 3500/10000 [1:53:06<2:26:15, 1.35s/it]
```

```
Epoch 3500: MSE loss = 123.61683654785156
Epoch 3500: Energy loss = 10.593358993530273
Epoch 3500: f1 loss = 0.01946951448917389
Epoch 3500: f2 loss = 0.10517560690641403
```



159.782883: 36%|


| 3600/10000 [1:56:29<2:34:28, 1.45s/it]

Epoch 3600: MSE loss = 120.44023895263672

Epoch 3600: Energy loss = 13.642826080322266

Epoch 3600: f1 loss = 11.408793449401855

Epoch 3600: f2 loss = 14.291030883789062



124.012421: 37%|


| 3700/10000 [1:59:53<2:40:21, 1.53s/it]

Epoch 3700: MSE loss = 108.57910919189453

Epoch 3700: Energy loss = 15.23764705657959

Epoch 3700: f1 loss = 0.1879320740699768

Epoch 3700: f2 loss = 0.007730198558419943




124.911423: 38%|




| 3800/10000 [2:03:21<2:22:06, 1.38s/it]

Epoch 3800: MSE loss = 108.26319885253906
 Epoch 3800: Energy loss = 16.00284194946289
 Epoch 3800: f1 loss = 0.10270579159259796
 Epoch 3800: f2 loss = 0.542682409286499




123.245995: 39%|



| 3900/10000 [2:06:49<2:33:45, 1.51s/it]

Epoch 3900: MSE loss = 108.59770965576172
 Epoch 3900: Energy loss = 14.602813720703125
 Epoch 3900: f1 loss = 0.0022757297847419977
 Epoch 3900: f2 loss = 0.04319242760539055




119.962463: 40%|



| 4000/10000 [2:10:18<2:35:05, 1.55s/it]

Epoch 4000: MSE loss = 104.53607177734375
 Epoch 4000: Energy loss = 15.408740997314453
 Epoch 4000: f1 loss = 0.012941724620759487
 Epoch 4000: f2 loss = 0.004710355773568153




56.942379: 41%|



| 4100/10000 [2:13:46<2:14:24, 1.37s/it]

Epoch 4100: MSE loss = 44.5427360534668
 Epoch 4100: Energy loss = 12.300982475280762
 Epoch 4100: f1 loss = 0.0912981852889061
 Epoch 4100: f2 loss = 0.007359548006206751



181.086777: 42%|


| 4200/10000 [2:17:11<2:28:31, 1.54s/it]

Epoch 4200: MSE loss = 39.45859146118164

Epoch 4200: Energy loss = 12.357421875

Epoch 4200: f1 loss = 128.32054138183594

Epoch 4200: f2 loss = 0.9502289891242981



81.503967: 43%|


| 4300/10000 [2:20:39<2:30:09, 1.58s/it]

Epoch 4300: MSE loss = 70.4593276977539

Epoch 4300: Energy loss = 11.044417381286621

Epoch 4300: f1 loss = 0.00013098414638079703


Epoch 4300: f2 loss = 8.919739775592461e-05



49.869175: 44%|

| 4400/10000 [2:24:01<2:07:05, 1.36s/it]


Epoch 4400: MSE loss = 33.705116271972656
 Epoch 4400: Energy loss = 15.837178230285645
 Epoch 4400: f1 loss = 0.3191782832145691
 Epoch 4400: f2 loss = 0.007703568786382675



59.592987: 45%|

| 4500/10000 [2:27:28<2:29:18, 1.63s/it]


Epoch 4500: MSE loss = 45.57386779785156
 Epoch 4500: Energy loss = 13.584615707397461
 Epoch 4500: f1 loss = 0.0007912294240668416
 Epoch 4500: f2 loss = 0.4337102472782135



144.371002: 46%|

| 4600/10000 [2:30:53<2:23:04, 1.59s/it]


Epoch 4600: MSE loss = 91.05794525146484
 Epoch 4600: Energy loss = 12.067108154296875
 Epoch 4600: f1 loss = 0.07221687585115433
 Epoch 4600: f2 loss = 41.17372512817383



133.939484: 47%|

| 4700/10000 [2:34:17<1:59:26, 1.35s/it]

Epoch 4700: MSE loss = 120.99455261230469
 Epoch 4700: Energy loss = 12.811744689941406
 Epoch 4700: f1 loss = 0.04469681158661842
 Epoch 4700: f2 loss = 0.08848919719457626



102.258751: 48%|


| 4800/10000 [2:37:43<2:15:22, 1.56s/it]

Epoch 4800: MSE loss = 85.24440002441406

Epoch 4800: Energy loss = 13.454198837280273

Epoch 4800: f1 loss = 3.0868983268737793

Epoch 4800: f2 loss = 0.47325554490089417



113.419037: 49%|


| 4900/10000 [2:41:10<2:10:04, 1.53s/it]

Epoch 4900: MSE loss = 101.9107666015625

Epoch 4900: Energy loss = 11.457845687866211

Epoch 4900: f1 loss = 0.008396034128963947

Epoch 4900: f2 loss = 0.04202989861369133



105.511551: 50%|




| 5000/10000 [2:44:35<1:48:17, 1.30s/it]

Epoch 5000: MSE loss = 95.16519165039062

Epoch 5000: Energy loss = 10.162062644958496

Epoch 5000: f1 loss = 0.17048360407352448

Epoch 5000: f2 loss = 0.013813060708343983



84.429184: 51%|




| 5100/10000 [2:47:59<2:02:04, 1.49s/it]

Epoch 5100: MSE loss = 72.90951538085938

Epoch 5100: Energy loss = 11.423993110656738

Epoch 5100: f1 loss = 0.005810747388750315

Epoch 5100: f2 loss = 0.08986225724220276




450.419678: 52%|




| 5200/10000 [2:51:21<2:03:37, 1.55s/it]

Epoch 5200: MSE loss = 83.13867950439453
 Epoch 5200: Energy loss = 15.673971176147461
 Epoch 5200: f1 loss = 346.8845520019531
 Epoch 5200: f2 loss = 4.722484588623047




73.158928: 53%|



| 5300/10000 [2:54:50<1:42:45, 1.31s/it]

Epoch 5300: MSE loss = 49.130828857421875
 Epoch 5300: Energy loss = 12.511900901794434
 Epoch 5300: f1 loss = 11.332438468933105
 Epoch 5300: f2 loss = 0.18375878036022186




74.434235: 54%|



| 5400/10000 [2:58:15<1:57:11, 1.53s/it]

Epoch 5400: MSE loss = 61.00910186767578
 Epoch 5400: Energy loss = 13.421124458312988
 Epoch 5400: f1 loss = 0.0034101279452443123
 Epoch 5400: f2 loss = 0.0005997715052217245




74.866058: 55%|



| 5500/10000 [3:01:44<1:42:53, 1.37s/it]

Epoch 5500: MSE loss = 59.64122009277344
 Epoch 5500: Energy loss = 15.090228080749512
 Epoch 5500: f1 loss = 0.018885118886828423
 Epoch 5500: f2 loss = 0.11572907865047455




61.153816: 56%|



| 5600/10000 [3:05:14<1:53:24, 1.55s/it]

Epoch 5600: MSE loss = 46.63825988769531
 Epoch 5600: Energy loss = 14.51231575012207
 Epoch 5600: f1 loss = 0.0013059302000328898
 Epoch 5600: f2 loss = 0.0019316618563607335




96.455498: 57%|



| 5700/10000 [3:09:03<2:02:19, 1.71s/it]

Epoch 5700: MSE loss = 80.50399017333984
 Epoch 5700: Energy loss = 14.689935684204102
 Epoch 5700: f1 loss = 1.241753339767456
 Epoch 5700: f2 loss = 0.019820403307676315




74.350777: 58%|



| 5800/10000 [3:12:52<1:49:28, 1.56s/it]

Epoch 5800: MSE loss = 62.22893142700195
 Epoch 5800: Energy loss = 12.084227561950684
 Epoch 5800: f1 loss = 0.002362170722335577
 Epoch 5800: f2 loss = 0.0352545902132988




170.655182: 59%|



| 5900/10000 [3:16:28<1:51:50, 1.64s/it]

Epoch 5900: MSE loss = 158.05934143066406
 Epoch 5900: Energy loss = 12.592269897460938
 Epoch 5900: f1 loss = 0.00320369191467762
 Epoch 5900: f2 loss = 0.00037174910539761186




128.249451: 60%|



| 6000/10000 [3:20:04<1:39:52, 1.50s/it]

Epoch 6000: MSE loss = 90.55164337158203
 Epoch 6000: Energy loss = 14.795979499816895
 Epoch 6000: f1 loss = 0.005473766475915909
 Epoch 6000: f2 loss = 22.8963623046875




79.735710: 61%|



| 6100/10000 [3:24:19<1:37:32, 1.50s/it]

Epoch 6100: MSE loss = 69.58786010742188
 Epoch 6100: Energy loss = 9.80348014831543
 Epoch 6100: f1 loss = 0.26468169689178467
 Epoch 6100: f2 loss = 0.07968537509441376




62.122410: 62%|



| 6200/10000 [3:27:48<1:23:25, 1.32s/it]

Epoch 6200: MSE loss = 46.6884880065918
 Epoch 6200: Energy loss = 15.059398651123047
 Epoch 6200: f1 loss = 0.00489977328106761
 Epoch 6200: f2 loss = 0.3696241080760956




104.171524: 63%|




| 6300/10000 [3:31:22<1:37:51, 1.59s/it]

Epoch 6300: MSE loss = 91.8114013671875
 Epoch 6300: Energy loss = 12.31852912902832
 Epoch 6300: f1 loss = 0.014882151037454605
 Epoch 6300: f2 loss = 0.02671211212873459




85.573715: 64%|



| 6400/10000 [3:34:49<1:19:10, 1.32s/it]

Epoch 6400: MSE loss = 74.81857299804688
 Epoch 6400: Energy loss = 10.572466850280762
 Epoch 6400: f1 loss = 0.1481216847896576
 Epoch 6400: f2 loss = 0.03455916419625282




94.458603: 65%|



| 6500/10000 [3:38:17<1:28:41, 1.52s/it]

Epoch 6500: MSE loss = 76.0320053100586
 Epoch 6500: Energy loss = 18.423580169677734
 Epoch 6500: f1 loss = 0.0023540270049124956
 Epoch 6500: f2 loss = 0.0006665781838819385




98.334946: 66%|



| 6600/10000 [3:41:51<1:17:36, 1.37s/it]

Epoch 6600: MSE loss = 83.62423706054688
 Epoch 6600: Energy loss = 14.708995819091797
 Epoch 6600: f1 loss = 0.00014611922961194068
 Epoch 6600: f2 loss = 0.0015667061088606715




111.897476: 67%|



| 6700/10000 [3:45:24<1:26:50, 1.58s/it]

Epoch 6700: MSE loss = 99.60868835449219
 Epoch 6700: Energy loss = 12.274126052856445
 Epoch 6700: f1 loss = 0.012867055833339691
 Epoch 6700: f2 loss = 0.0017963218269869685



112.661697: 68%|




| 6800/10000 [3:48:52<1:21:45, 1.53s/it]

Epoch 6800: MSE loss = 98.41017150878906

Epoch 6800: Energy loss = 14.093432426452637

Epoch 6800: f1 loss = 0.15528617799282074

Epoch 6800: f2 loss = 0.002811260288581252



137.543533: 69%|




| 6900/10000 [3:52:26<1:18:06, 1.51s/it]

Epoch 6900: MSE loss = 125.50194549560547

Epoch 6900: Energy loss = 12.027569770812988

Epoch 6900: f1 loss = 0.002276299288496375

Epoch 6900: f2 loss = 0.01174819003790617




84.038971: 70%|



| 7000/10000 [3:55:53<1:14:34, 1.49s/it]

Epoch 7000: MSE loss = 71.02739715576172
 Epoch 7000: Energy loss = 13.001675605773926
 Epoch 7000: f1 loss = 0.004543713293969631
 Epoch 7000: f2 loss = 0.005355107132345438




112.081505: 71%|



| 7100/10000 [3:59:27<1:08:14, 1.41s/it]

Epoch 7100: MSE loss = 99.32754516601562
 Epoch 7100: Energy loss = 12.460679054260254
 Epoch 7100: f1 loss = 0.06804623454809189
 Epoch 7100: f2 loss = 0.22523410618305206



93.930611: 72%|




| 7200/10000 [4:02:54<1:11:09, 1.52s/it]

Epoch 7200: MSE loss = 82.81610107421875

Epoch 7200: Energy loss = 11.084904670715332

Epoch 7200: f1 loss = 0.004419948905706406

Epoch 7200: f2 loss = 0.025184322148561478



95.196014: 73%|




| 7300/10000 [4:06:22<1:13:55, 1.64s/it]

Epoch 7300: MSE loss = 85.10733795166016

Epoch 7300: Energy loss = 10.070757865905762

Epoch 7300: f1 loss = 0.015638114884495735

Epoch 7300: f2 loss = 0.0022822923492640257




66.426727: 74%|



| 7400/10000 [4:09:48<58:31, 1.35s/it]

Epoch 7400: MSE loss = 54.1306037902832
 Epoch 7400: Energy loss = 12.259466171264648
 Epoch 7400: f1 loss = 0.031549349427223206
 Epoch 7400: f2 loss = 0.005110486410558224




74.118683: 75%|



| 7500/10000 [4:13:19<1:06:28, 1.60s/it]

Epoch 7500: MSE loss = 58.947784423828125
 Epoch 7500: Energy loss = 15.165406227111816
 Epoch 7500: f1 loss = 0.00507731270045042
 Epoch 7500: f2 loss = 0.0004151529283262789



112.079765: 76%|




| 7600/10000 [4:16:51<59:31, 1.49s/it]

Epoch 7600: MSE loss = 95.60213470458984

Epoch 7600: Energy loss = 14.913339614868164

Epoch 7600: f1 loss = 0.007361915893852711

Epoch 7600: f2 loss = 1.5569281578063965



109.589630: 77%|




| 7700/10000 [4:20:25<53:59, 1.41s/it]

Epoch 7700: MSE loss = 91.1583023071289

Epoch 7700: Energy loss = 17.238521575927734

Epoch 7700: f1 loss = 0.9923180937767029

Epoch 7700: f2 loss = 0.2004840224981308




60.587887: 78%|



| 7800/10000 [4:24:00<58:03, 1.58s/it]

Epoch 7800: MSE loss = 46.44302749633789
 Epoch 7800: Energy loss = 14.135884284973145
 Epoch 7800: f1 loss = 0.005710158962756395
 Epoch 7800: f2 loss = 0.0032651035580784082




67.913170: 79%|



| 7900/10000 [4:27:30<46:37, 1.33s/it]

Epoch 7900: MSE loss = 57.100833892822266
 Epoch 7900: Energy loss = 10.788277626037598
 Epoch 7900: f1 loss = 0.014254715293645859
 Epoch 7900: f2 loss = 0.009805936366319656




118.308685: 80%|



| 8000/10000 [4:30:57<53:11, 1.60s/it]

Epoch 8000: MSE loss = 105.23762512207031
 Epoch 8000: Energy loss = 10.98361587524414
 Epoch 8000: f1 loss = 0.36695030331611633
 Epoch 8000: f2 loss = 1.7204970121383667




54.546375: 81%|



| 8100/10000 [4:34:25<42:51, 1.35s/it]

Epoch 8100: MSE loss = 39.528953552246094
 Epoch 8100: Energy loss = 13.3958101272583
 Epoch 8100: f1 loss = 0.1911027580499649
 Epoch 8100: f2 loss = 1.4305086135864258




126.357414: 82%|



| 8200/10000 [4:37:56<47:55, 1.60s/it]

Epoch 8200: MSE loss = 75.69306945800781
 Epoch 8200: Energy loss = 12.745185852050781
 Epoch 8200: f1 loss = 37.862144470214844
 Epoch 8200: f2 loss = 0.0570170059800148




64.284576: 83%|




| 8300/10000 [4:41:29<38:21, 1.35s/it]

Epoch 8300: MSE loss = 52.57515335083008
 Epoch 8300: Energy loss = 11.286636352539062
 Epoch 8300: f1 loss = 0.3183002471923828
 Epoch 8300: f2 loss = 0.10448973625898361




181.053802: 84%|




| 8400/10000 [4:44:59<41:46, 1.57s/it]

Epoch 8400: MSE loss = 131.85736083984375
 Epoch 8400: Energy loss = 12.667200088500977
 Epoch 8400: f1 loss = 36.328189849853516
 Epoch 8400: f2 loss = 0.20104777812957764




53.667122: 85%|



| 8500/10000 [4:48:29<33:42, 1.35s/it]

Epoch 8500: MSE loss = 39.792659759521484
 Epoch 8500: Energy loss = 13.852118492126465
 Epoch 8500: f1 loss = 0.0017169727943837643
 Epoch 8500: f2 loss = 0.020625099539756775



73.689117: 86%|




| 8600/10000 [4:52:00<35:48, 1.53s/it]

Epoch 8600: MSE loss = 56.43167495727539

Epoch 8600: Energy loss = 17.253765106201172

Epoch 8600: f1 loss = 0.003165320260450244

Epoch 8600: f2 loss = 0.000515670282766223



41.980648: 87%|




| 8700/10000 [4:55:33<29:27, 1.36s/it]

Epoch 8700: MSE loss = 27.354936599731445

Epoch 8700: Energy loss = 14.616411209106445

Epoch 8700: f1 loss = 0.0050104218535125256

Epoch 8700: f2 loss = 0.004290134180337191



525402.000000: 88%|




| 8800/10000 [4:59:08<30:24, 1.52s/it]

Epoch 8800: MSE loss = 71.76712799072266

Epoch 8800: Energy loss = 15.027170181274414

Epoch 8800: f1 loss = 525296.375

Epoch 8800: f2 loss = 18.85719108581543



91.817207: 89%|




| 8900/10000 [5:02:41<29:17, 1.60s/it]

Epoch 8900: MSE loss = 77.6731948852539

Epoch 8900: Energy loss = 13.827415466308594

Epoch 8900: f1 loss = 0.2539275586605072

Epoch 8900: f2 loss = 0.06266950815916061




87.104935: 90%|



| 9000/10000 [5:06:15<22:44, 1.36s/it]

Epoch 9000: MSE loss = 71.15874481201172
 Epoch 9000: Energy loss = 15.872835159301758
 Epoch 9000: f1 loss = 0.05758344382047653
 Epoch 9000: f2 loss = 0.01577364094555378




62.515560: 91%|



| 9100/10000 [5:09:52<23:21, 1.56s/it]


Epoch 9100: MSE loss = 47.712852478027344
 Epoch 9100: Energy loss = 13.635639190673828
 Epoch 9100: f1 loss = 0.0022993297316133976
 Epoch 9100: f2 loss = 1.164768099784851



100.165710: 92%|

| 9200/10000 [5:13:20<19:44, 1.48s/it]


Epoch 9200: MSE loss = 88.60223388671875
 Epoch 9200: Energy loss = 11.088390350341797
 Epoch 9200: f1 loss = 0.46549370884895325
 Epoch 9200: f2 loss = 0.009596948511898518



137.139374: 93%|

| 9300/10000 [5:16:49<18:14, 1.56s/it]

Epoch 9300: MSE loss = 118.29869079589844
 Epoch 9300: Energy loss = 18.438114166259766
 Epoch 9300: f1 loss = 0.05490834638476372
 Epoch 9300: f2 loss = 0.34766116738319397




65.959839: 94%|



| 9400/10000 [5:20:20<15:20, 1.53s/it]

Epoch 9400: MSE loss = 50.20254898071289
 Epoch 9400: Energy loss = 14.71146011352539
 Epoch 9400: f1 loss = 0.510416567325592
 Epoch 9400: f2 loss = 0.5354181528091431




58.800320: 95%|



| 9500/10000 [5:23:56<11:16, 1.35s/it]

Epoch 9500: MSE loss = 41.593177795410156
 Epoch 9500: Energy loss = 16.789154052734375
 Epoch 9500: f1 loss = 0.4135289788246155
 Epoch 9500: f2 loss = 0.004458288662135601




97.049278: 96%|



| 9600/10000 [5:27:37<09:41, 1.45s/it]

Epoch 9600: MSE loss = 84.96973419189453
 Epoch 9600: Energy loss = 12.059536933898926
 Epoch 9600: f1 loss = 0.006253623869270086
 Epoch 9600: f2 loss = 0.013752649538218975




45.455379: 97%|



| 9700/10000 [5:31:05<07:40, 1.54s/it]


Epoch 9700: MSE loss = 29.070720672607422
 Epoch 9700: Energy loss = 13.887406349182129
 Epoch 9700: f1 loss = 2.4899983406066895
 Epoch 9700: f2 loss = 0.007254406344145536



709.324646: 98%|

| 9800/10000 [5:34:41<04:29, 1.35s/it]


Epoch 9800: MSE loss = 123.96302032470703
 Epoch 9800: Energy loss = 12.735799789428711
 Epoch 9800: f1 loss = 534.519775390625
 Epoch 9800: f2 loss = 38.10610580444336



113.975517: 99%|

| 9900/10000 [5:38:15<02:38, 1.59s/it]

Epoch 9900: MSE loss = 72.50479125976562
 Epoch 9900: Energy loss = 13.458986282348633
 Epoch 9900: f1 loss = 27.332685470581055
 Epoch 9900: f2 loss = 0.6790546774864197



91.352242: 100%|

| 10000/10000 [5:41:46<00:00, 2.05s/it]

```
print(s3_faces.shape)
print(s1_faces_tmp.shape)
print(s1_faces.shape)

(970, 3, 3)
(1078, 3, 3)
(970, 3, 3)

lf1, lf2 = s1_faces_tmp, s3_faces

centroids_lf1 = np.array([compute_centroid(face) for face in lf1])
centroids_lf2 = np.array([compute_centroid(face) for face in lf2])

closest_faces = []
closest_faces_idx = []

for centroid in centroids_lf2:
    distances = np.linalg.norm(centroids_lf1 - centroid, axis=1)
    closest_face_idx = np.argmin(distances)
    while closest_face_idx in closest_faces_idx:
        distances[closest_face_idx] = max(distances) + 1
        # print(closest_face_idx, np.argmin(distances))
        closest_face_idx = np.argmin(distances)
    closest_faces.append(lf1[closest_face_idx])
    closest_faces_idx.append(closest_face_idx)
print(len(closest_faces_idx))
print(max(closest_faces_idx))

# for centroid in centroids_lf2:
#     distances = np.linalg.norm(centroids_lf1 - centroid, axis=1)
#     closest_face_idx = np.argmin(distances)
#     while closest_face_idx in closest_faces_idx:
#         distances[closest_face_idx] = max(distances) + 1
#         # print(closest_face_idx, np.argmin(distances))
```

```

#         closest_face_idx = np.argmin(distances)
#         closest_faces.append(lf1[closest_face_idx])
#         closest_faces_idx.append(closest_face_idx)
# print(len(closest_faces_idx))
# print(max(closest_faces_idx))

970
1077

v1s = my_batch.pos.detach().numpy()[0]
v3s = s3.features.detach().squeeze(0).numpy()[0]
print(v1s, v3s)

[[ 3.1602659e+00 -1.0895985e+01  1.0343000e-02]
 [ 3.7360101e+00 -1.1155475e+01  4.6799000e-02]
 [ 5.8451301e-01 -1.9746441e+01  2.1269619e+01]
 ...
 [-1.0321780e+00  1.3775474e+01  8.3145781e+00]
 [-7.1533501e-01  1.1227485e+01  1.3906136e+01]
 [-3.3355400e-01  1.4022309e+01  7.2379370e+00]] [-1.09389 -
3.4720056  7.0445876]

np.save('shape1_baseline.npy', v1_remapped)
np.save('shape2_reconstruction.npy',
s3.features.detach().squeeze(0).numpy())

```