

Geometric Computer Vision: Final Project

This notebook contains a revisited implementation of the paper: **Shape Non-rigid Kinematics (SNK): A Zero-Shot Method for Non-Rigid Shape Matching via Unsupervised Functional Map Regularized Reconstruction** by Attaiki and Ovsjanikov (2024).

Imports

```
from pyFM.functional import FunctionalMapping
import numpy as np
import trimesh
from pyFM.mesh import TriMesh
import torch
from torch_geometric.data import Batch
from diffusion_net import DiffusionData, DiffusionOperatorsTransform,
DiffusionNet
from prism_decoder import PrismDecoder
import torch.optim as optim
from tqdm import tqdm
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d.art3d import Poly3DCollection
from matplotlib import cm
from loss import PrismRegularizationLoss
```

Load meshes

Let's first load the meshes of two dogs modified (non-rigid deformation using Blender):

```
# load some meshes
obj1 = r"./Samples/dog_small.obj"
obj2 = r"./Samples/dog_small_after_torture.obj"
mesh1, mesh2 = TriMesh(obj1), TriMesh(obj2)
```

Compute correspondance from mesh2 to mesh1 using functional maps

I chose the pyFM package investigated in HW3 to compute a correspondance map between the meshes, using WKS descriptors:

```
process_params = {
    'n_ev': (70, 70), # Number of eigenvalues on source and Target
    'subsample_step': 20, # In order not to use too many descriptors
    'descr_type': 'WKS', # WKS or HKS
}
model = FunctionalMapping(mesh1, mesh2)
```

```

model.preprocess(**process_params, verbose=False)
fit_params = {
    'w_descr': 1e0,
    'w_lap': 1e-2,
    'w_dcomm': 1e-1,
    'w_orient': 0
}
model.fit(**fit_params, verbose=False)
p2p_21 = model.get_p2p(n_jobs=1)

```

Visualization using meshplot package

```

import meshplot as mp

def double_plot_mp(myMesh1, myMesh2, cmap1=None, cmap2=None):
    p = mp.plot(myMesh1.vertlist, myMesh1.facelist, c=cmap1)
    v2_translated = myMesh2.vertlist.copy()
    f2_translated = myMesh2.facelist.copy()
    myMesh2_translated = TriMesh(v2_translated, f2_translated)
    myMesh2_translated.vertlist[:,0] =
myMesh2_translated.vertlist[:,0] + 30
    p.add_mesh(myMesh2_translated.vertlist,
myMesh2_translated.facelist, c=cmap2)

def visu(vertices):
    min_coord, max_coord =
np.min(vertices, axis=0, keepdims=True), np.max(vertices, axis=0, keepdims=
True)
    cmap = (vertices - min_coord) / (max_coord - min_coord)
    return cmap

cmap1 = visu(mesh1.vertlist); cmap2 = cmap1[p2p_21]
double_plot_mp(mesh1, mesh2, cmap1, cmap2)

{"model_id": "fc1770e0779940718bef04bff3f18226", "version_major": 2, "vers
ion_minor": 0}

```

Compute latent vector

Following section 4.2 in the paper, we compute a latent vector (of length d_1) per feature using DiffusionNet++ module, then apply max pooling to get a vector l of length d_1 .

Then we aim to reconstruct the first shape by concatenating l to each feature of the second mesh and passing those to the decoder.

```

def compute_l(my_batch, diffusion_net):
    my_batch.x = my_batch.pos.clone() # set the input features to the
positions
    output = diffusion_net(my_batch)

```

```

    l = output.x.max(dim=0).values
    return l

def reconstruct_s2(my_batch2, v2_t, l, decoder):
    l_expanded = l.unsqueeze(0).repeat(v2_t.shape[0],1)
    my_batch2.x = torch.cat((v2_t,l_expanded),dim=1) #.unsqueeze(0)
    # my_batch2.pos = my_batch2.pos.unsqueeze(0)
#my_batch2.pos.clone().clone() # set the input features to the
positions
    s3 = decoder(my_batch2)
    return s3

```

Prepare first mesh for compatibility with DiffusionNet++ model

```

mesh1_diff = trimesh.load(obj1)
v1, f1 = np.array(mesh1.vertices), np.array(mesh1.faces)
v1_t = torch.from_numpy(v1)
f1_t = torch.from_numpy(f1)
data1 = DiffusionData(pos=v1_t, face=f1_t.T)
diffusion_transform = DiffusionOperatorsTransform(n_eig=50) #97
compute the diffusion net operators with 97 eigenvalues
data1 = diffusion_transform(data1)
my_batch = Batch.from_data_list([data1])

```

Prepare second mesh for compatibility with DiffusionNet++ model

```

mesh2_diff = trimesh.load(obj2)
v2, f2 = np.array(mesh2.vertices), np.array(mesh2.faces)
data2 = DiffusionData(pos=torch.from_numpy(v2),
face=torch.from_numpy(f2).T)
diffusion_transform = DiffusionOperatorsTransform(n_eig=50) #97
compute the diffusion net operators with 97 eigenvalues
data2 = diffusion_transform(data2)
my_batch2 = Batch.from_data_list([data2])
my_batch2.pos = my_batch2.pos.unsqueeze(0)
#####
v2_t = torch.Tensor(v2)

```

Make parameters differentiable - for backward propagation

```

my_batch.pos.requires_grad = True
v2_t.requires_grad = True
my_batch2.pos.requires_grad = True

```

Prepare visualization functions

```

def double_plot(myMesh1, myMesh2, cmap1=None, cmap2=None,
title1='Mesh1',title2='Mesh2'):

```

```

fig = plt.figure(figsize=plt.figaspect(0.5))
ax1 = fig.add_subplot(1, 2, 1, projection='3d')
ax1.plot_trisurf(myMesh1.vertlist[:, 0], myMesh1.vertlist[:, 1],
myMesh1.vertlist[:, 2],
                 triangles=myMesh1.facelist, cmap='viridis',
facecolors=cmap1)
ax1.set_title(title1)
ax2 = fig.add_subplot(1, 2, 2, projection='3d')
ax2.plot_trisurf(myMesh2.vertlist[:, 0], myMesh2.vertlist[:, 1],
myMesh2.vertlist[:, 2],
                 triangles=myMesh2.facelist, cmap='viridis',
facecolors=cmap2)
ax2.set_title(title2)

def visu_face_colors(s1_faces_raw, s1_faces, s2_faces, s3_faces):
    cmap = cm.get_cmap('viridis', s3_faces.shape[0]) # Use a colormap
    with 500 distinct colors
    face_colors = cmap(np.linspace(0, 1, s3_faces.shape[0]))

    average_z = np.mean(s2_faces[:, :, 2], axis=1)
    sorted_indices = np.argsort(average_z)

    sorted_s3_faces = s3_faces[sorted_indices]
    sorted_s2_faces = s2_faces[sorted_indices]
    sorted_s1_faces = s1_faces[sorted_indices]
    fig = plt.figure()
    ax1 = fig.add_subplot(131, projection='3d')
    ax2 = fig.add_subplot(132, projection='3d')
    ax3 = fig.add_subplot(133, projection='3d')
    for i, face in enumerate(s3_faces):
        # poly3d = [[vertices[vert_idx] for vert_idx in face]]
        poly3d_1 = [sorted_s1_faces[i]]
        poly3d_2 = [sorted_s2_faces[i]]
        poly3d_1_raw = [s1_faces_raw[i]]
        poly3d_3 = [sorted_s3_faces[i]]
        ax1.add_collection3d(Poly3DCollection(poly3d_1,
facecolors=face_colors[i], linewidths=1, edgecolors=None, alpha=.8))
        ax1.add_collection3d(Poly3DCollection(poly3d_1_raw,
facecolors='b', linewidths=1, edgecolors=None, alpha=.1))
        ax2.add_collection3d(Poly3DCollection(poly3d_2,
facecolors=face_colors[i], linewidths=1, edgecolors=None, alpha=.8))
        ax3.add_collection3d(Poly3DCollection(poly3d_3,
facecolors=face_colors[i], linewidths=1, edgecolors=None, alpha=.8))
        ax1.set_xlim([s1_faces[:, :, 0].min(), s1_faces[:, :, 0].max()])
        ax1.set_ylim([s1_faces[:, :, 1].min(), s1_faces[:, :, 1].max()])
        ax1.set_zlim([s1_faces[:, :, 2].min(), s1_faces[:, :, 2].max()])
        ax2.set_xlim([s2_faces[:, :, 0].min(), s2_faces[:, :, 0].max()])
        ax2.set_ylim([s2_faces[:, :, 1].min(), s2_faces[:, :, 1].max()])
        ax2.set_zlim([s2_faces[:, :, 2].min(), s2_faces[:, :, 2].max()])

```

```

ax3.set_xlim([s3_faces[:, :, 0].min(), s3_faces[:, :, 0].max()])
ax3.set_ylim([s3_faces[:, :, 1].min(), s3_faces[:, :, 1].max()])
ax3.set_zlim([s3_faces[:, :, 2].min(), s3_faces[:, :, 2].max()])
plt.show()

```

Handling face matching for visualization

```

def compute_centroid(face):
    return np.mean(face, axis=0)

def find_closest_faces(lf1, lf2):
    centroids_lf1 = np.array([compute_centroid(face) for face in lf1])
    centroids_lf2 = np.array([compute_centroid(face) for face in lf2])

    closest_faces = []
    closest_faces_idx = []

    for centroid in centroids_lf2:
        distances = np.linalg.norm(centroids_lf1 - centroid, axis=1)
        closest_face_idx = np.argmin(distances)
        while closest_face_idx in closest_faces_idx:
            distances[closest_face_idx] = max(distances) + 1
            closest_face_idx = np.argmin(distances)
        closest_faces.append(lf1[closest_face_idx])
        closest_faces_idx.append(closest_face_idx)
    return closest_faces_idx

vls = my_batch.pos.reshape(-1, 3)
fls = my_batch.face.t()
sl_faces_tmp = vls[fls].detach().numpy()

```

Initialization of the models (DiffusionNet, PrismDecoder)

```

L_SPACE_SIZE = 42
diffusion_net = DiffusionNet(3, L_SPACE_SIZE)
decoder = PrismDecoder(v1.shape[1]+L_SPACE_SIZE, v1.shape[0])

```

Training process

```

# %debug
lambda_E = 30
losses_MSE = []
losses_energy = []
losses = []
get_energy_loss = PrismRegularizationLoss(100)
optimizer_diffnet = optim.AdamW(diffusion_net.parameters(), lr=0.001,
weight_decay=0.001)
optimizer_decoder = optim.AdamW(decoder.parameters(), lr=0.001,
weight_decay=0.001)

```

```

for epoch in (pbar:=tqdm(range(2000))):
    l = compute_l(my_batch, diffusion_net)
    l.retain_grad()
    s3 = reconstruct_s2(my_batch2, v2_t, l, decoder)
    s3.features.retain_grad()
    v1_remapped = v1[p2p_21]
    loss =
torch.nn.functional.mse_loss(torch.Tensor(v1_remapped), s3.features)
    loss_E = get_energy_loss(s3.transformed_prism, s3.rotations,
s3.pos.reshape(-1, 3), s3.face)
    total_loss = loss + loss_E*lambda_E
    losses_MSE.append(loss.detach())
    losses_energy.append(loss_E.detach())
    losses.append(total_loss.detach())
    pbar.set_description(f"{loss.detach():3f}")
    total_loss.retain_grad()
    total_loss.backward()
    # import ipdb; ipdb.set_trace()
    optimizer_diffnet.step()
    optimizer_diffnet.zero_grad()
    optimizer_decoder.step()
    optimizer_decoder.zero_grad()
    if (epoch%20 == 0 and epoch<=101) or epoch%100 == 0:
        print(f'Epoch {epoch}: MSE loss = {loss}')
        print(f'Epoch {epoch}: Energy loss = {loss_E}')
        verts = my_batch2.pos.reshape(-1, 3)
        faces = my_batch2.face.t()
        s2_faces = verts[faces].detach().numpy()
        s3_faces = s3.transformed_prism.detach().squeeze(0).numpy()
        closest_faces_idx = find_closest_faces(s1_faces_tmp, s3_faces)
#find_closest_faces(s3_faces, s1_faces_tmp)
        s1_faces = s1_faces_tmp[closest_faces_idx]
        visu_face_colors(s1_faces_tmp, s1_faces, s2_faces, s3_faces)
        plt.show()

```

0%|

```

| 0/2000 [00:00<?, ?it/s]C:\Users\Hadassa-Port\AppData\Local\Temp\
ipykernel_12508\321993680.py:15: UserWarning: Using a target size
(torch.Size([1, 487, 3])) that is different to the input size
(torch.Size([487, 3])). This will likely lead to incorrect results due
to broadcasting. Please ensure they have the same size.

```

loss =

```

torch.nn.functional.mse_loss(torch.Tensor(v1_remapped), s3.features)
C:\Users\Hadassa-Port\Desktop\hadassa\Toar 2\Semester 1\Geometric
Computer Vision\SNK\loss.py:28: UserWarning: Using torch.cross without
specifying the dim arg is deprecated.

```


Please either pass the dim explicitly or simply use
torch.linalg.cross.

The default value of dim will change to agree with that of
linalg.cross in a future release. (Triggered internally at ..\aten\

```
src\ATen\native\Cross.cpp:66.)
  normal = torch.cross(edge1, edge2)
380.823242:  0%|
| 0/2000 [00:00<?, ?it/s]
```


Epoch 0: MSE loss = 380.8232421875
 Epoch 0: Energy loss = 55.170536041259766

C:\Users\Hadassa-Port\AppData\Local\Temp\ipykernel_12508\2963957283.py:13: MatplotlibDeprecationWarning: The get_cmap function was deprecated in Matplotlib 3.7 and will be removed two minor releases later. Use ``matplotlib.colormaps[name]`` or ``matplotlib.colormaps.get_cmap(obj)`` instead.
 cmap = cm.get_cmap('viridis', s3_faces.shape[0]) # Use a colormap with 500 distinct colors




```
37.519310:  1%|
| 20/2000 [00:48<33:17,  1.01s/it]
```

Epoch 20: MSE loss = 37.519309997558594
 Epoch 20: Energy loss = 0.06417781859636307




```
18.741827:  2%|
| 40/2000 [01:34<47:36,  1.46s/it]
```


Epoch 40: MSE loss = 18.7418270111084
Epoch 40: Energy loss = 0.002584312343969941




16.039995: 3%| ██████████
| 60/2000 [02:10<32:58, 1.02s/it]

Epoch 60: MSE loss = 16.039995193481445
Epoch 60: Energy loss = 0.00724763935431838




15.875034: 4%| ██████████
| 80/2000 [02:40<32:13, 1.01s/it]

Epoch 80: MSE loss = 15.87503433227539
Epoch 80: Energy loss = 0.006988008506596088




15.714131: 5%|██████████
| 100/2000 [03:11<32:33, 1.03s/it]

Epoch 100: MSE loss = 15.714131355285645
Epoch 100: Energy loss = 0.006078184582293034



14.708385: 10%|██████████
| 200/2000 [05:04<31:01, 1.03s/it]


Epoch 200: MSE loss = 14.708385467529297
Epoch 200: Energy loss = 0.01923668012022972



14.443155: 15%|██████████
| 300/2000 [07:04<29:41, 1.05s/it]

Epoch 300: MSE loss = 14.443155288696289
Epoch 300: Energy loss = 0.02276437170803547

Epoch 600: MSE loss = 13.603837966918945
Epoch 600: Energy loss = 0.032199397683143616




13.452585: 35%|



| 700/2000 [14:36<22:22, 1.03s/it]

Epoch 700: MSE loss = 13.452585220336914
Epoch 700: Energy loss = 0.03310574218630791




12.586988: 40%|




| 800/2000 [16:26<22:13, 1.11s/it]

Epoch 800: MSE loss = 12.58698844909668
Epoch 800: Energy loss = 0.0529869943857193




12.083807: 45%|




| 900/2000 [18:20<19:17, 1.05s/it]

Epoch 900: MSE loss = 12.083806991577148

Epoch 900: Energy loss = 0.05478142201900482




11.784428: 50%|



| 1000/2000 [20:11<16:32, 1.01it/s]

Epoch 1000: MSE loss = 11.784427642822266

Epoch 1000: Energy loss = 0.05672043561935425




11.695845: 55%|




| 1100/2000 [22:26<20:13, 1.35s/it]

Epoch 1100: MSE loss = 11.695844650268555

Epoch 1100: Energy loss = 0.05680251121520996




11.427905: 60%|



| 1200/2000 [24:51<24:02, 1.80s/it]

Epoch 1200: MSE loss = 11.427905082702637

Epoch 1200: Energy loss = 0.06552743166685104




11.417137: 65%|



| 1300/2000 [27:07<11:57, 1.03s/it]

Epoch 1300: MSE loss = 11.417137145996094

Epoch 1300: Energy loss = 0.0595039464533329




11.263186: 70%|



| 1400/2000 [28:56<10:09, 1.02s/it]

Epoch 1400: MSE loss = 11.263185501098633

Epoch 1400: Energy loss = 0.06322402507066727




11.132992: 75%|



| 1500/2000 [30:49<08:45, 1.05s/it]

Epoch 1500: MSE loss = 11.132991790771484

Epoch 1500: Energy loss = 0.06326287239789963




11.285318: 80%|



| 1600/2000 [32:40<06:39, 1.00it/s]

Epoch 1600: MSE loss = 11.285318374633789

Epoch 1600: Energy loss = 0.056058600544929504




11.038894: 85%|




| 1700/2000 [34:30<04:51, 1.03it/s]

Epoch 1700: MSE loss = 11.038893699645996

Epoch 1700: Energy loss = 0.05881454795598984




10.967781: 90%|



| 1800/2000 [36:20<03:22, 1.01s/it]

Epoch 1800: MSE loss = 10.967781066894531

Epoch 1800: Energy loss = 0.062212999910116196




10.702518: 95%|



| 1900/2000 [38:13<01:37, 1.03it/s]

Epoch 1900: MSE loss = 10.702518463134766

Epoch 1900: Energy loss = 0.06658665835857391



10.562922: 100%|



| 2000/2000 [40:08<00:00, 1.20s/it]

Plot the loss functions

```
# Create the figure and subplots
fig, axs = plt.subplots(2, 2, figsize=(18, 5))
l_energy = [lambda_E*a for a in losses_energy]
# Plot MSE loss
axs[0,0].plot(range(len(losses_MSE)), losses_MSE, label='Loss MSE',
marker='o', color='b')
axs[0,0].set_xlabel('Epoch'); axs[0,0].set_ylabel('Loss MSE')
axs[0,0].set_title('Loss MSE per epoch'); axs[0,0].legend();
axs[0,0].grid(True)
```

```

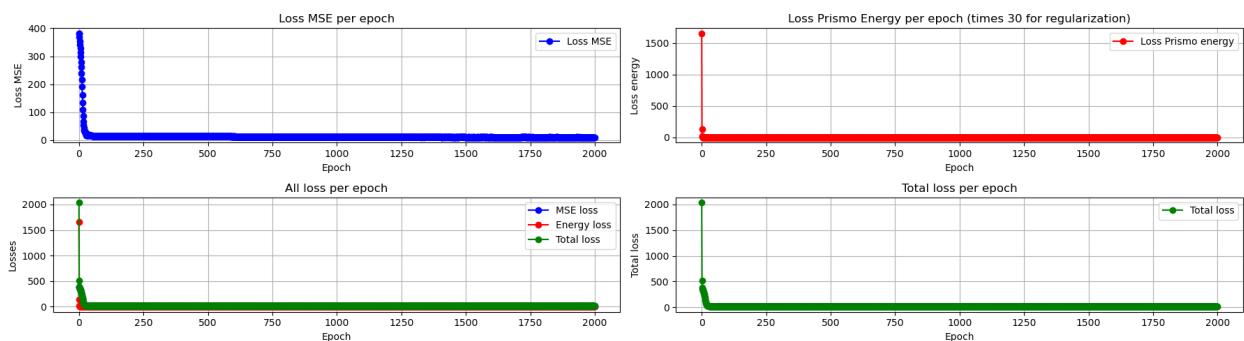
# Plot Energy loss
axs[0,1].plot(range(len(losses_energy)), l_energy, label='Loss Prismo
energy', marker='o', color='r')
axs[0,1].set_xlabel('Epoch'); axs[0,1].set_ylabel('Loss energy')
axs[0,1].set_title(f'Loss Prismo Energy per epoch (times {lambda_E}
for regularization)'); axs[0,1].legend(); axs[0,1].grid(True)

# Plot total loss
axs[1,0].plot(range(len(losses_MSE)), losses_MSE, label='MSE loss',
marker='o', color='b')
axs[1,0].plot(range(len(l_energy)), l_energy, label='Energy loss',
marker='o', color='r')
axs[1,0].plot(range(len(losses)), losses, label='Total loss',
marker='o', color='g')
axs[1,0].set_xlabel('Epoch'); axs[1,0].set_ylabel('Losses')
axs[1,0].set_title('All loss per epoch'); axs[1,0].legend();
axs[1,0].grid(True)

# Plot total loss
axs[1,1].plot(range(len(losses)), losses, label='Total loss',
marker='o', color='g')
axs[1,1].set_xlabel('Epoch'); axs[1,1].set_ylabel('Total loss')
axs[1,1].set_title('Total loss per epoch'); axs[1,1].legend();
axs[1,1].grid(True)

# Show the plot
plt.tight_layout()
plt.show()

```



Save first shape (remapped based on functional map) and its reconstruction.

```

np.save('shape1_baseline.npy', v1_remapped)
np.save('shape2_reconstruction.npy',
s3.features.detach().squeeze(0).numpy())

```